

### III. Model documentation and write-up

Information included in this section may be shared publicly with challenge results. You can respond to these questions in an e-mail or as an attached file. Please number your responses.

1. Who are you (mini-bio) and what do you do professionally? If you are on a team, please complete this block for each member of the team.

In my other life, I worked various roles as part of a data analytics team - as a systems admin, DBA, data warehouse architect and reports developer. I now work as a freelance data scientist. Over the past 2 years I've been studying machine learning through data competitions, solving all sorts of fascinating challenges.

2. What motivated you to compete in this challenge?

Satellites can 'see' the air that we breathe? It's ingenious. And I love a good challenge. Whenever I come across such a challenge, I'm reminded of the capacity of machine learning to alleviate the broader issues facing our society today. We can build solutions that affect millions of people, hopefully change their lives for the better.

3. High level summary of your approach: what did you do and why?

My approach was almost entirely data-centric. Initially I experimented with the original competition dataset and a variety of models. I quickly hit a plateau where local CV and LB scores didn't improve.

I shifted my focus to the external datasets, and settled for a simple gradient boosting model to speed up my experiments. I used the MAIAC satellite data and GFS forecasts for the PM2.5 model. For the NO2 model, I used OMI and TROPOMI satellite data, GFS forecasts and GEOS-CF NO2 hindcasts.

Choice of GFS variables was based on relevant literature and availability of data throughout the training and testing period i.e. Jan 2017 to Aug 2021. Generally, I selected parameters affecting or similar to air humidity, soil temperature, soil humidity, air temperature, wind velocity, wind direction and rainfall/ precipitation.

GFS forecasts from upto 3 days preceding the date of interest seemed to improve the predictions, so I created multiple datasets with different lookback periods. I also realized that separate models for each location performed better than a single model. Since the test set includes dataset from the past, I decided to use a k-fold (without shuffle) cross validation strategy instead of time-based splits. The final solution is an average ensemble of 45 models (3 datasets x 5 folds x 3 locations) for the PM2.5 model and 30 models (2 datasets x 5 folds x 3 locations) for the NO2 model.

4. Do you have any useful charts, graphs, or visualizations from the process?  
No.

5. Copy and paste the 3 most impactful parts of your code and explain what each does and how it helped your model.

```
##impute missing MAIAC data by expanding the region of interest with  
small increments. The same idea is applied for TROPOMI data  
bounds=row.geometry.bounds  
for OFF in [i*0.01 for i in range(1,5)]:  
    #increment area with a small offset  
    ixs=np.where((lon>=bounds[0]-OFF) & (lon<=bounds[2]+OFF) &  
(lat>=bounds[1]-OFF) & (lat<=bounds[3]+OFF))  
    d_mean=np.nanmean(data[ixs])  
    d_median=np.nanmedian(data[ixs])  
    if not np.isnan(d_mean):  
        break
```

6. Please provide the machine specs and time you used to run your model.

- CPU (model): Intel Xeon @ 2.30GHz, 4vCPUs
- GPU (model or N/A): N/A
- Memory (GB): 32
- OS: Ubuntu 21.04
- Train duration: 2 hours (see *runtime.html* for details)
- Inference duration: 10 minutes

7. Anything we should watch out for or be aware of in using your model (e.g. code quirks, memory requirements, numerical stability issues, etc.)?

Downloading GFS data is the biggest bottleneck in the pipeline. The NCAR data archive server (<https://rda.ucar.edu>) has a limit of 10 concurrent requests, and seems to process only 2 requests at a time while the rest are queued. Try using separate NCAR accounts if running the NO2 and PM2.5 preprocessing scripts concurrently.

It's also worth noting that requests for GFS data sometimes fail for no good reason. Retrying the failed request usually works, so my current solution is a loop which could potentially run forever :(.

Refer to *src/data/extract\_gfs.py*

```
while True:  
    REQUESTS,PARAM_PATHS = generate_requests()  
    results = pqdm(REQUESTS, download_data,  
n_jobs=MAX_CONCURRENT_REQUESTS,argument_type=None)  
    completed = [res for res in results if res['success']==True]  
    if len(completed)>0:  
        print('Completed requests: ',completed)  
    failed = [res for res in results if res['success']==False]  
    if len(failed)>0:  
        print('The following requests failed. Retrying...')  
        print(failed)  
    else:  
        break
```

8. Did you use any tools for data preparation or exploratory data analysis that aren't listed in your code submission?

No

9. How did you evaluate the performance of the model other than the provided metric, if at all?

I optimized RMSE for the PM2.5 model and Huber loss for the NO2 model. I tracked the best R2 validation score for both models.

10. What are some other things you tried that didn't necessarily make it into the final workflow (quick overview)?

RNNs/ Transformers with a rolling time window, but sequential data didn't seem to improve forecasts. I also briefly experimented with segmentation models but with the low spatial resolution and missing data it was a long shot to begin with.

11. If you were to continue working on this problem for the next year, what methods or techniques might you try in order to build on your work so far? Are there other fields or features you felt would have been very helpful to have?

Feature selection, tuning - I ran out of time before I could identify the most useful features and properly tune the model parameters before my last submission. With around 300 features it's likely that the models could be improved or at least simplified with comprehensive feature selection.

Higher resolution data - I didn't get a chance to explore the ECMWF dataset, which has a higher spatial resolution than GFS. Integrating the dataset in the model could improve forecasts.

Redundancy - Sporadic outages of the [GEOS-CF data server](#) got me thinking about the resilience of this solution in a production environment due to its dependence on external datasets. I would consider building different models that work with different combinations of the available datasets and at different scales e.g. using 3,6,12,24 hour forecasts.